

# Event processing with Kafka for the Pythonista

Pradeep Gowda

pradeep@btbytes.com


October 11, 2016

Indypy

# Survey

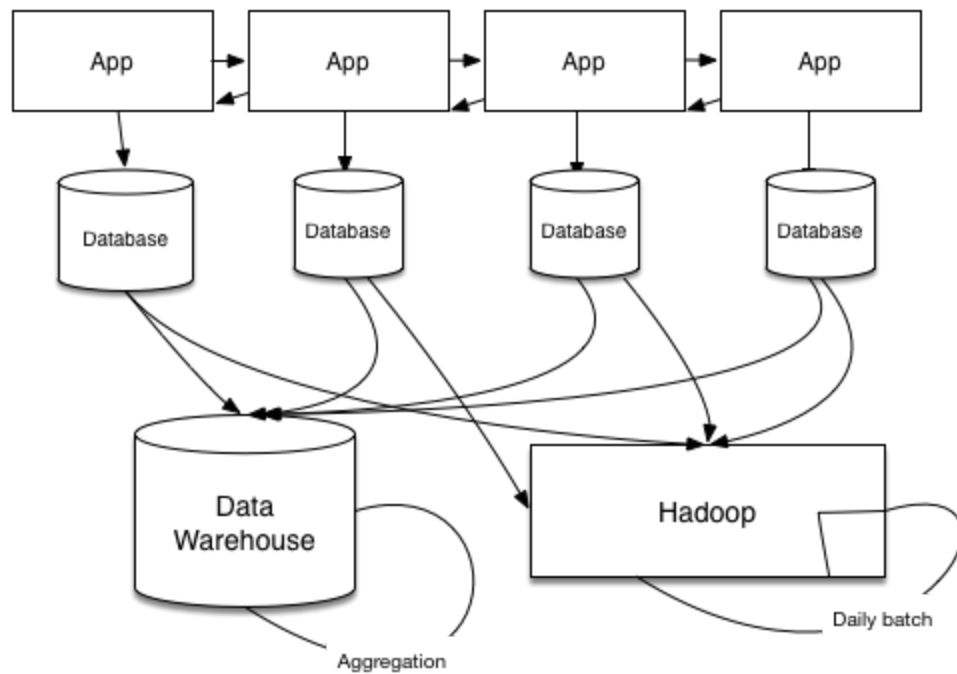
# **The nature of data**

## "pre-streaming" data architecture




(3-layers, n-layers)

# Problem: Pipeline sprawl



# Problem: Everything is synchronous



## How did we end up here?

- multiple downstream users
- evolution of uses over time not anticipated at design time
- previous solutions to address this -- data lake, OLAP

SAME  
DATA

in

DIFFERENT  
FORM

Denormalisation

Indexes

Caching

Aggregations



**All your data is event streams**

# Event data example

1. User views a product

```
{  
  "product_id": 1234,  
  "time": "2016-10-05T12:12:12",  
  "user_id": 3298,  
}
```

## Event data example (continued..)

2. User checks-in to a new location
3. User catches a poke-thing
4. Log file

```
64.242.88.10 - - [07/Mar/2004:17:29:59 -0800] "GET /twiki,  
64.242.88.10 - - [07/Mar/2004:17:31:39 -0800] "GET /twiki,  
64.242.88.10 - - [07/Mar/2004:17:35:35 -0800] "GET /twiki,
```

5. Sensors

# What are tables?

Tables

key1	val3
key2	val2

# Tables over time

## Tables over time

time = 0	<table><tr><td>key1</td><td>val1</td></tr></table>	key1	val1	PUT(key1, val1)		
key1	val1					
time = 1	<table><tr><td>key1</td><td>val1</td></tr><tr><td>key2</td><td>val2</td></tr></table>	key1	val1	key2	val2	PUT(key2, val2)
key1	val1					
key2	val2					
time = 2	<table><tr><td>key1</td><td>val3</td></tr><tr><td>key2</td><td>val2</td></tr></table>	key1	val3	key2	val2	PUT(key1, val3)
key1	val3					
key2	val2					

# Sounds like ...



## *Write-Ahead Logging (WAL)*

- ❖ The Write-Ahead Logging Protocol:
  - ① Must **force** the **log record** for an update *before* the corresponding **data page** gets to disk.
  - ② Must **write all log records** for a Xact *before commit*.
- ❖ #1 guarantees Atomicity.
- ❖ #2 guarantees Durability.
  
- ❖ Exactly how is logging (and recovery!) done?
  - We'll study the ARIES algorithms.

## WAL-first rule

---

- The log record for an operation is always written to disk before the affected data pages
  - WAL first rule!
- In case of a crash, the WAL is replayed to reconstruct the unsaved changes

Write-Ahead Logging is critical for:

- Durability
  - Once you commit, your data is safe
- Consistency
  - No corruption on crash



## Advanced features

---



But there's more!

The Write-Ahead Log *also* allows:

- Online backup
- Point-in-time Recovery
- Replication

# **What is Kafka?**

Kafka is a modern distributed Platform for data streams

# Three properties of a streaming platform

ie., Kafka

1. Publish and subscribe to stream of records (similar to Queues)
2. Store stream of records in a fault-tolerant way
3. Process streams of records as they occur

# What is it used for?

Build real time

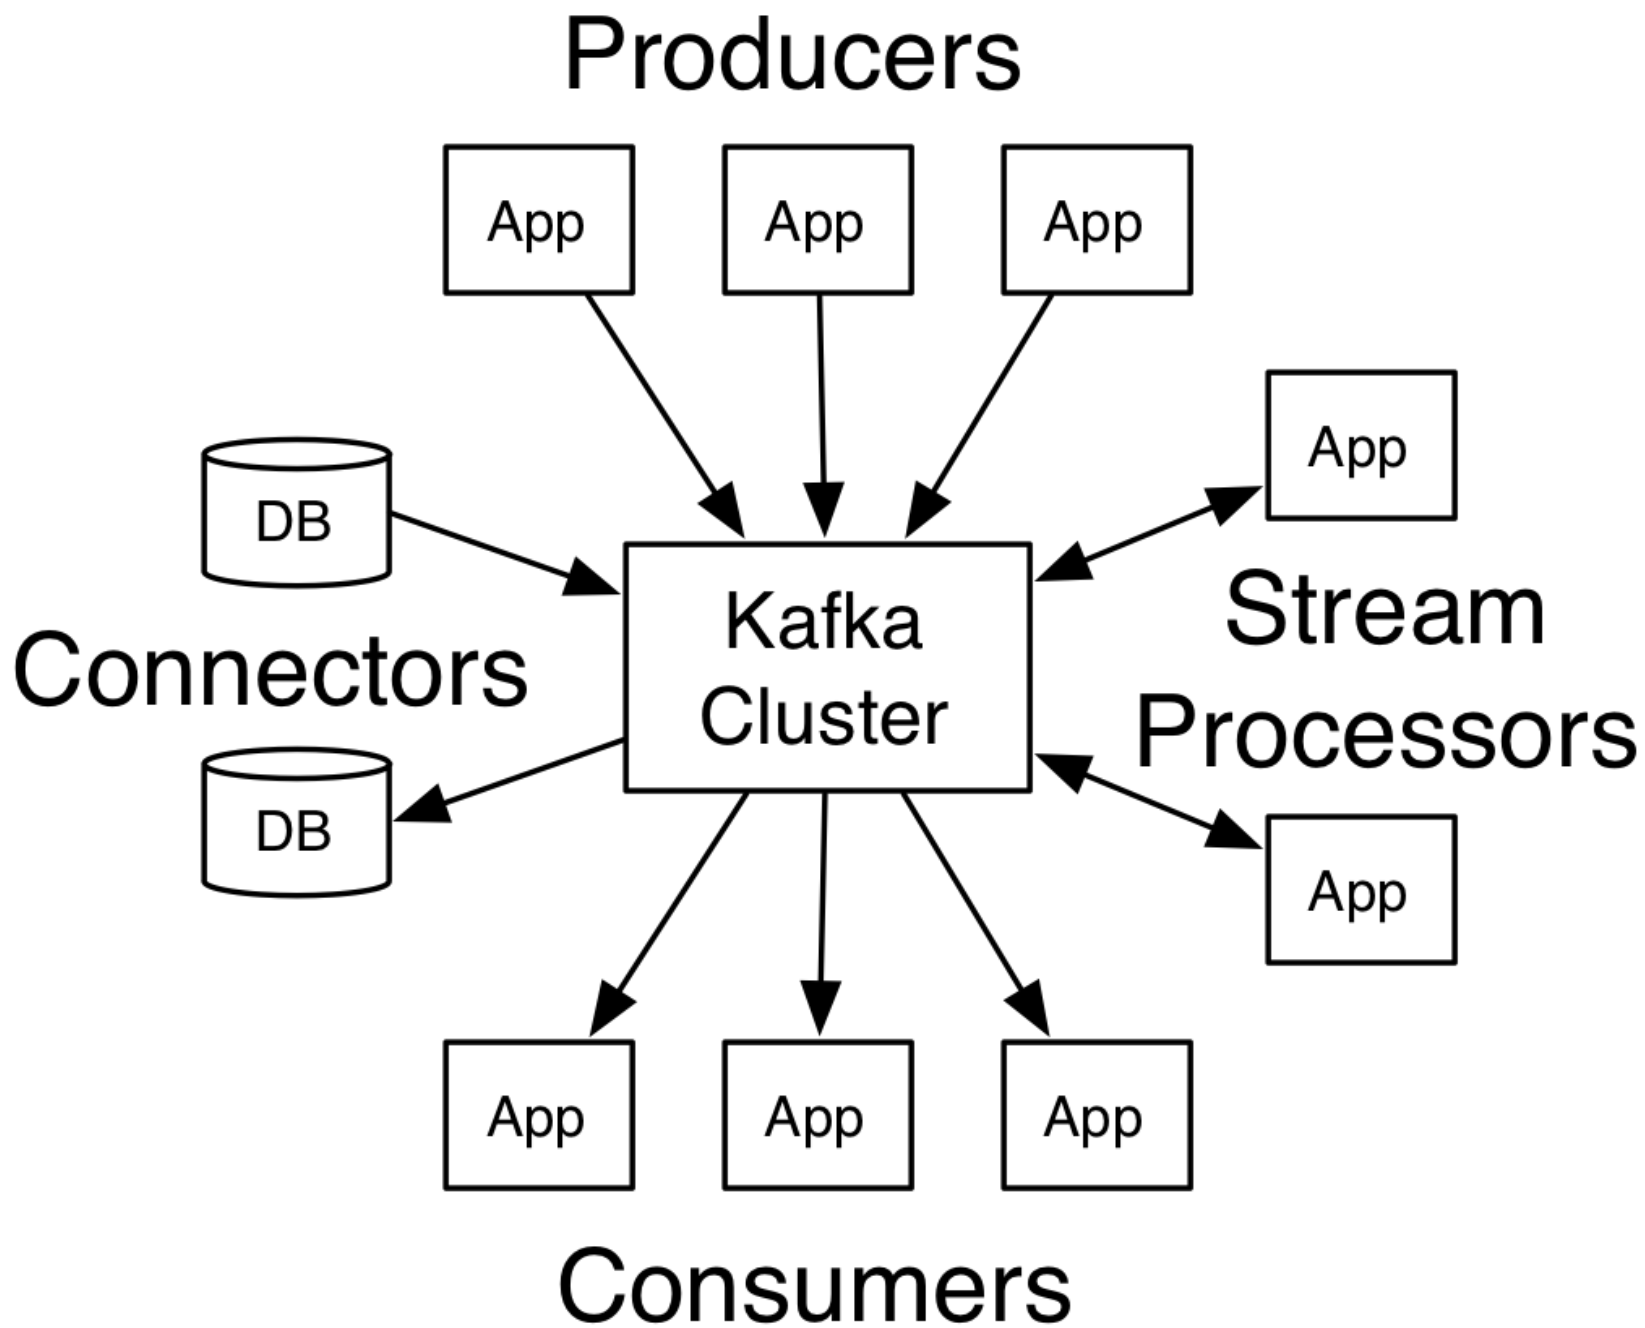
- streaming pipelines to ferry data between systems/apps
- streaming applications to transform/react to stream of data

# Kafka concepts

- Kafka is run as cluster of servers
- Kafka cluster stores streams of *records* in categories called *topics*
- Each record has this form — ( `key` , `value` , `timestamp` )


# Kafka APIs

1. Producer API
2. Consumer API
3. Streams API
4. Connector API



# Streaming Data

INSTEAD, EMBRACE THE LOG






# Kafka Topic

- example of a topic: `poke-thing-capture`
- every topic is multi-subscriber. `0-n`

## Anatomy of a Topic

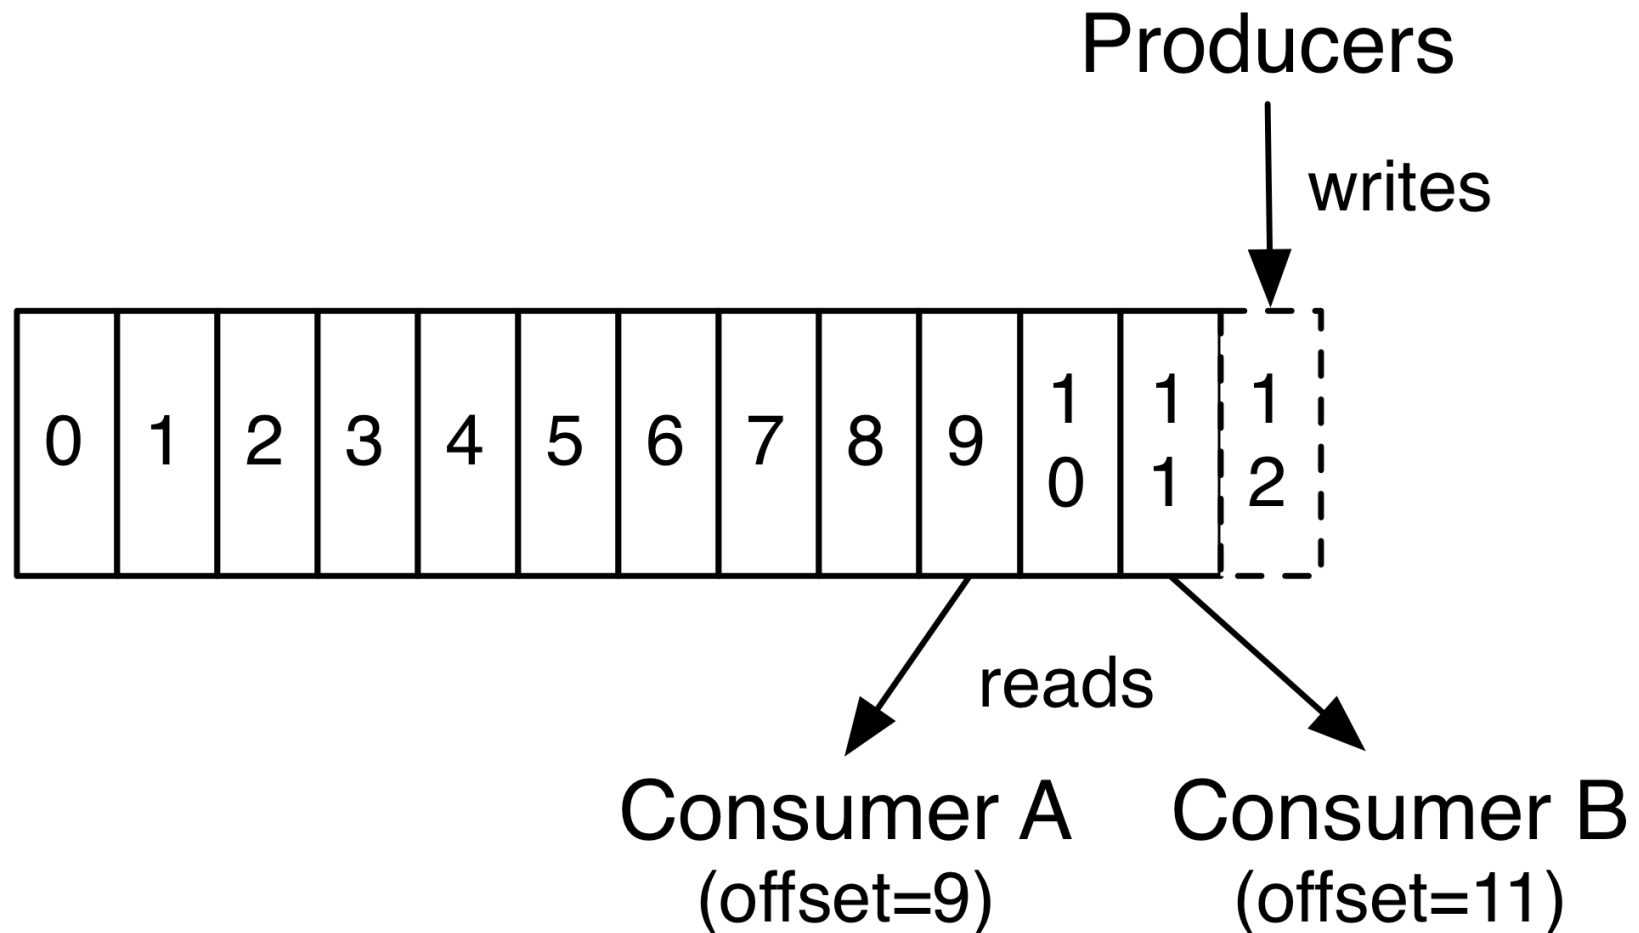


- Partition = immutable sequence of records.
- one topic may be split into multiple partitions
- ... **Append only**
- retained for `x` period (default= `1d` )

# Producer

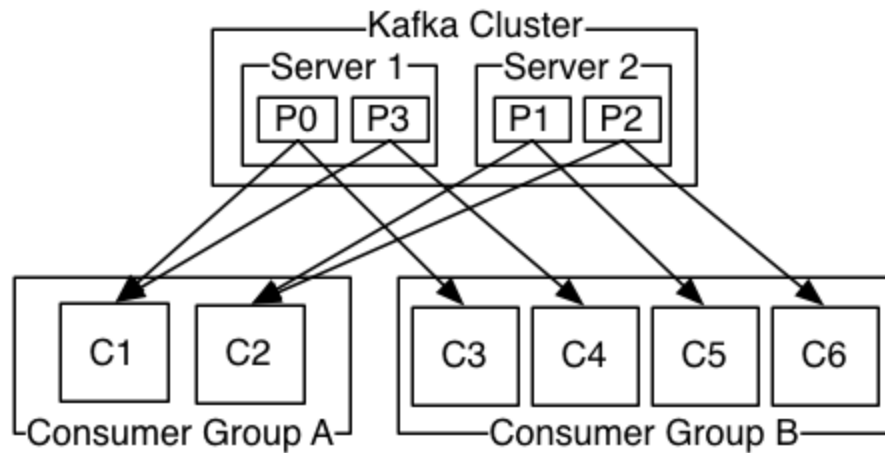
- Producers publish data to topics

# Consumers



- Every consumer has an "offset".
- This offset is controlled by the consumer

# Consumer Groups



- records load balanced over consumer instances

## Total ordering

- Kafka only provides a total order over records within a partition, not between different partitions in a topic.

# Distribution

- The partitions of the log are distributed over the servers in the Kafka cluster
- Each partition is replicated across a configurable number of servers for fault tolerance.
- Each partition has one server which acts as the "leader" and zero or more servers which act as "followers".

## Leaders and followers

- The leader handles all read and write requests for the partition while the followers passively replicate the leader.
- If the leader fails, one of the followers will automatically become the new leader.
- Each server acts as a leader for some of its partitions and a follower for others so load is well balanced within the cluster.
- c.f. **zookeeper**

# Kafka as ...

- Messaging system
  - queuing (-once you read the data, it's gone)
  - pub/sub
- Storage system
- Stream processing



## Use cases

- Messaging (À la RabbitMQ)
- Website activity tracking
- Metrics
- Log aggregation
- Stream processing
- [Event Sourcing](#)
- Commit log (c.f log compaction)

Kafka is a top level Apache project — <http://kafka.apache.org/>



HOME

INTRODUCTION

QUICKSTART

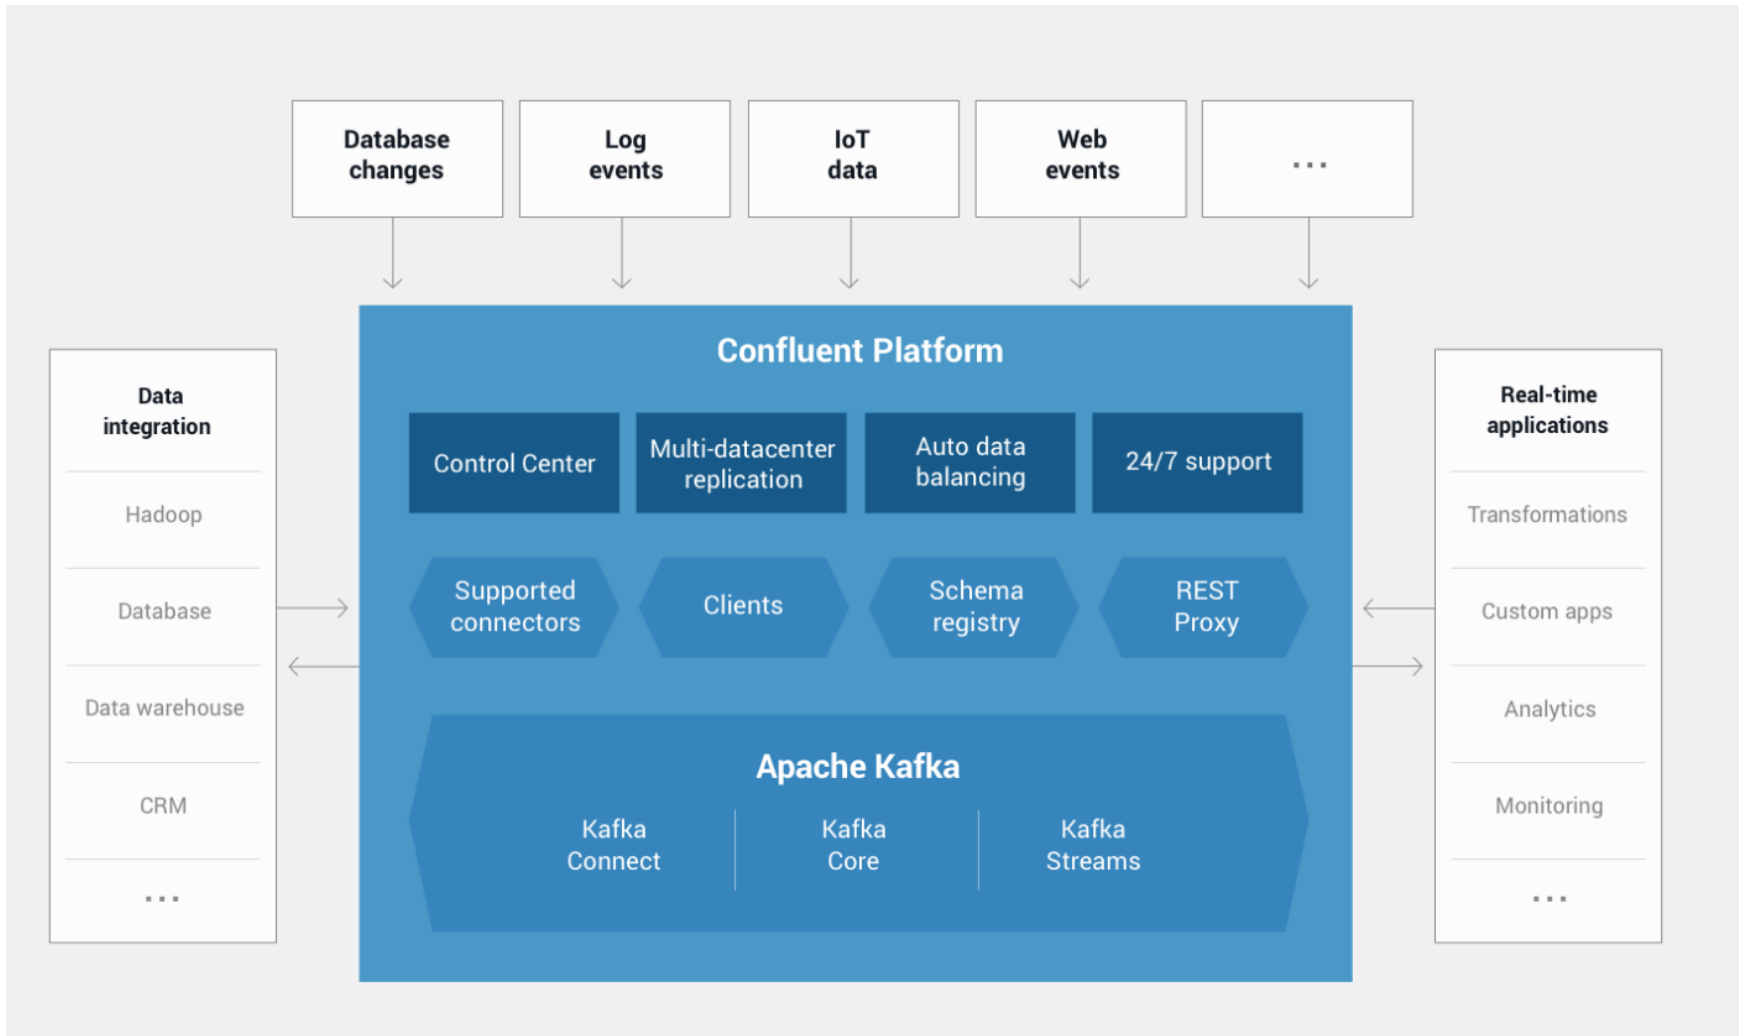
USE CASES



## PUBLISH & SUBSCRIBE

to streams of data like a  
messaging system

- Runs on the JVM
- Written in Scala
- [Apache Community Development - A Maturity Model for](#)



People:

- Jay Kreps
- Neha Narkhede

# Amazon Kinesis

Amazon Kinesis services make it easy to work with real-time streaming data in the AWS cloud.

Get started with Amazon Kinesis

- Firehose — load massive volumes of streaming data into AWS
- Analytics — analyze streaming data with standard SQL
- Streams — Build custom applications that process or analyze streaming data

## **Three key differences**

1. Scale
2. Guarantees
3. Stream Processing

## Scale

- Hundreds of MBs of Througput
- Store TBs of data
- Commodity Hardware
- $O(1)$  writes
- Scalability of a file system

## Guarantees

- Strict ordering (M1, M2)
- A consumer instance sees records in the order they are stored in the log.
- Handling server failures. (For a topic with replication factor  $N$ , we will tolerate up to  $N-1$  server failures without losing any records committed to the log.)

## **Distributed by design**

- Replication
- Fault tolerance
- Partitioning
- Elastic scaling



## Real life example

Linkedin

- >1.2 trillion messages per day written
- >3.4 trillion messages per day read
- 1PB of stored data

**Demo**

# Console demo

See `notes.txt`

# Python demo

```
pip install kafka-python
```

Output:

See [example-log.txt](#)

Thank you!

---

Image Credits —

- [@martinkl](#),
- [@jaykreps](#)
- [kafka project](#)
- [Readings in Database Systems, 3rd Ed](#)
- [Heikki Linnakangas](#)