

# Building IoT Data Pipelines With Python

Logan Wendholt


IoT Technical Lead, Bastian Solutions

# Overview


- IoT in 10 seconds
- Pipeline Basics
- Common Hardware Deployments
- Pipeline Design Patterns and Examples
- Cloud-based Pipelines / “Serverless”
- Hybrid Architectures

# Internet of Things

- Connect things... to the internet
  - Or maybe just your private network
- Why?
  - **DATA**





# What's in a Pipeline?



# Things to Consider


- What kind of data?
- How much?
- How fast?
- What type of “flow”?



- Hardware limitations?


# Things to Consider

A handful of devices in your apartment?




# Things to Consider

Hundreds of sensors in your factory?



# Things to Consider

Thousands of sensors deployed  
*anywhere?*





# Things to Consider

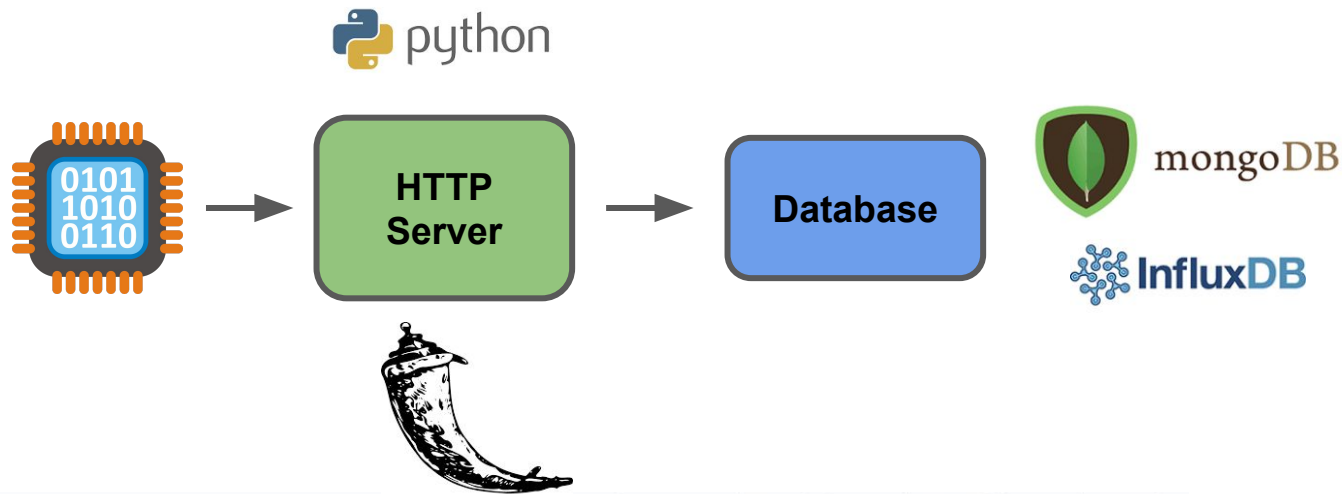
**How do I choose?**

# Client-Server Architecture

- May want to use a client-server pattern if...
  - Devices natively use HTTP or other client-server protocol
  - Data flow pattern is better suited for request-response format
    - Receive input -> send to cloud -> receive response -> perform action
  - Need to integrate with existing architecture
    - Can't spin up a new server, and access to database is restricted to a REST API?

# Client-Server Architecture

If you're using HTTP:



 pallets / flask

 Code

 Issues 37

 Pull requests 15

 Insights

# Flask

 Watch

2,222

 Star

39,290

 Fork

11,406

The Python micro framework for building web applications. <https://www.palletsprojects.com/p/flask/>

python

flask

wsgi


web-framework

werkzeug

jinja

# Client-Server Architecture

If not:



**Twisted**

[twisted / twisted](#)

[Code](#)

[Pull requests](#) 95

[Projects](#) 0

[Insights](#)

[Watch](#)

173

[Star](#)

2,953

[Fork](#)

801

Event-driven networking engine written in Python. <https://www.twistedmatrix.com>

[python](#) [twisted](#) [async](#) [async-python](#) [network](#) [event-driven](#) [http](#) [irc](#) [dns](#) [xmpp](#) [imap](#) [smtp](#) [tls](#) [ssl](#)

# Client-Server Architecture

- Can everything run on a single PC?




Local PC

**But... what if you need more resources?**

# Client-Server Architecture

If you need to scale:



# Client-Server Architecture

- Pros:
  - Familiar concepts and tools (web dev)
  - Can support HTTP via Flask or custom protocols via Twisted
- Cons:
  - Custom protocols can be problematic
  - HTTP can be inefficient for low-resource IoT devices
  - Need additional code to support real-time data streaming

**Need a different approach?**

# Publish-Subscribe Architecture


- May want to consider pub/sub if...
  - Sending data asynchronously
  - Small data packets (sensor data)
  - Power or bandwidth are a concern
- De facto IoT pub/sub standard: MQTT



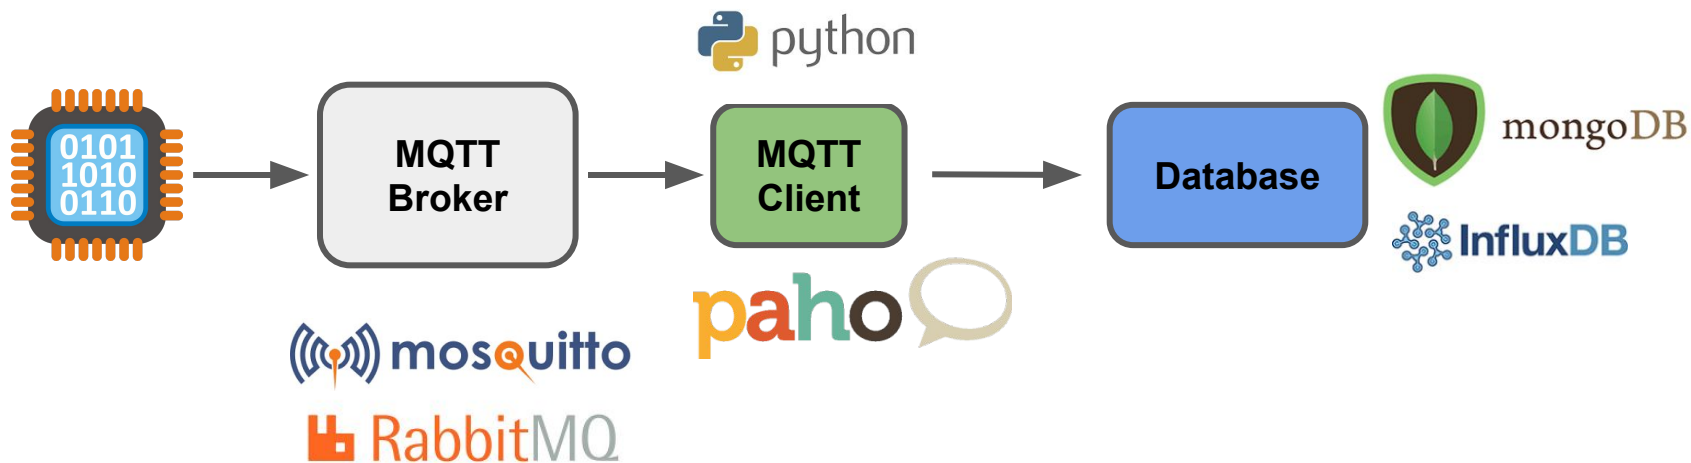


# What's MQTT?

- Message Queuing Telemetry Transport
  - Publish / subscribe messaging protocol
  - Application layer of network stack -- replaces HTTP
- Popular option for IoT devices
  - Very little data overhead
  - Well-suited to transmitting sensor data
    - Wake up
    - Publish
    - Go back to sleep

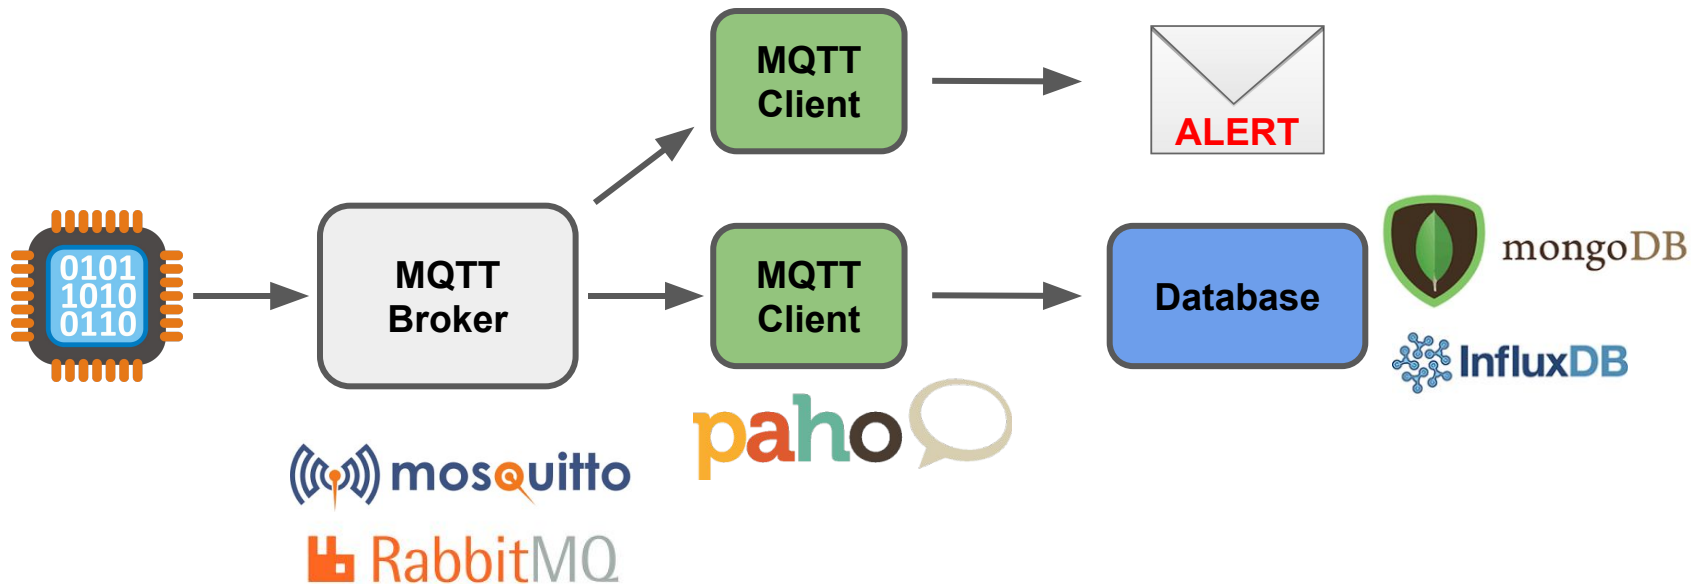


# Pub/Sub with MQTT



# Pub/Sub with MQTT

 python



# Pub/Sub with MQTT

- Pros:

- Pub/Sub makes it easy to stream data to multiple endpoints
  - Alerts system?
  - Real-time visualizations?
- Decouple data processing from message routing
  - Can swap out message broker if necessary

- Cons:

- Devices need to speak MQTT!
  - If you have to translate, it's probably not worth it
- Typically limited to one client working as “database adaptor”
  - Can lead to bottlenecks in high-throughput scenarios


# Pub/Sub with MQTT

- How to solve bottleneck / scaling issues?

**Use the cloud!**

# IoT + Cloud Pipeline

- Cloud-based, fully managed MQTT broker services



# IoT + Cloud Pipeline



- Function-as-a-Service (FaaS)
  - Self-contained functions (Python!)
  - Trigger based on events, inputs, etc.
  - Use cases:
    - ETL
    - Alerts
    - Analytics
    - Storage




Image Credit: [Austen Collins](#) - [ServerlessConf Austin '17](#)

- Cloud provider manages deployment, scaling, maintenance -- you just provide code!

# IoT + Cloud Pipeline

- Managed database solutions
  - Have the cloud provider run the database for you
  - Only charged for the storage and requests
  - Scale up or down fast!






# Serverless Architecture

Managed MQTT broker + FaaS + Managed Database = **Serverless**

- Exclusively use managed services rather than setting up VMs or servers
  - Let the cloud providers handle maintenance, installation, security updates, etc.
  - Scale up or down almost instantly
  - Only pay for what you use (requests, compute time, storage)
  - Focus development efforts where it counts: on the code!



# Serverless Architecture

Managed MQTT broker + FaaS + Managed Database = **Serverless**

- Infrastructure-as-Code
  - Create code templates that describe the resources you need
    - Databases, compute functions, other services
  - “Serverless Framework”
    - Platform-agnostic infrastructure as code
    - Easily deploy and scale pipelines
    - Less cloud provider lock-in



# Serverless Opportunities

- Easy to hook serverless pipeline into additional managed services
  - AWS IoT Analytics, Kinesis, Redshift
  - Google ML Engine, Dataflow, Bigquery
- Powerful analytics, machine learning, data warehousing options
- Beyond the scope of this talk

**Go explore!**

# Serverless Costs

## Scenario:

- Want to monitor temperature changes throughout a large warehouse
  - Place temperature sensors near all the doors
  - Send out a push alert to smartphone when temperature increases rapidly
- Want to use cloud solution for all processing and storage
  - How about AWS?
    - AWS IoT Core, Lambda, DynamoDB, Simple Notification Service



# Serverless Costs


## **Straight to the cloud:**

- 20 devices, report every 5 minutes: FREE
- 20 devices, report once per minute: \$1
- 20 devices, report every 30 seconds: \$3
- 20 devices, report once per second: \$110

**Lots of requests -> lots of \$\$\$**



# Edge Processing

- Process data locally -> only talk to cloud when necessary
  - Also works with unreliable network connection scenarios
  - AWS Greengrass, Google Cloud IoT Edge, Azure IoT Edge



# Back to the Scenario...

- Use AWS Greengrass to process data at edge
  - Local Lambda functions running on a Raspberry Pi -- still using Python!



- If temperature anomaly is detected, send request to cloud to send an alert
  - Store all temp data locally -- no need to stream to cloud
  - Maybe periodically send a backup of historical data?

# Serverless Costs

## **Straight to the cloud:**

- 20 devices, report every 5 minutes: FREE
- 20 devices, report once per minute: \$1
- 20 devices, report every 30 seconds: \$3
- 20 devices, report once per second: \$110



# Serverless Costs

~~Straight to the cloud:~~ **Edge + Cloud**

- 20 devices, report every 5 minutes: FREE
- 20 devices, report once per minute: ~~\$1~~ **FREE\***
- 20 devices, report every 30 seconds: ~~\$3~~ **FREE\***
- 20 devices, report once per second: ~~\$110~~ **FREE\***

\*Perform bulk of compute at edge and remain within free tier on requests

# Edge Processing Recap

- Dramatically reduce costs by doing pre-processing locally
- Only send data to cloud when necessary
- Great compromise
  - flexibility of the cloud
  - low operational costs of running on-prem
- Can still use Python functions as the primary method of computation

# To wrap it up...

- Wide range of possibilities for building data pipelines in Python
  - Flask / Twisted for client-server architectures
  - Paho MQTT client + RabbitMQ or Mosquitto for pub/sub architectures
  - Managed IoT services from Google, Microsoft, Amazon for cloud-based pub/sub
- Keep the cloud in mind
  - Serverless architectures can be very powerful!
  - Be careful with calculating your costs
  - Cloud-based pipeline with edge pre-processing: best of both worlds?
- Use the best tools for the job
  - Consider your use case, existing resources, and experience before deciding on an approach
  - Don't get cut by the cutting-edge!

Thanks!