

# **Jupyter / Docker Architecture**


**High Level Overview  
of what we are doing**

# Topics

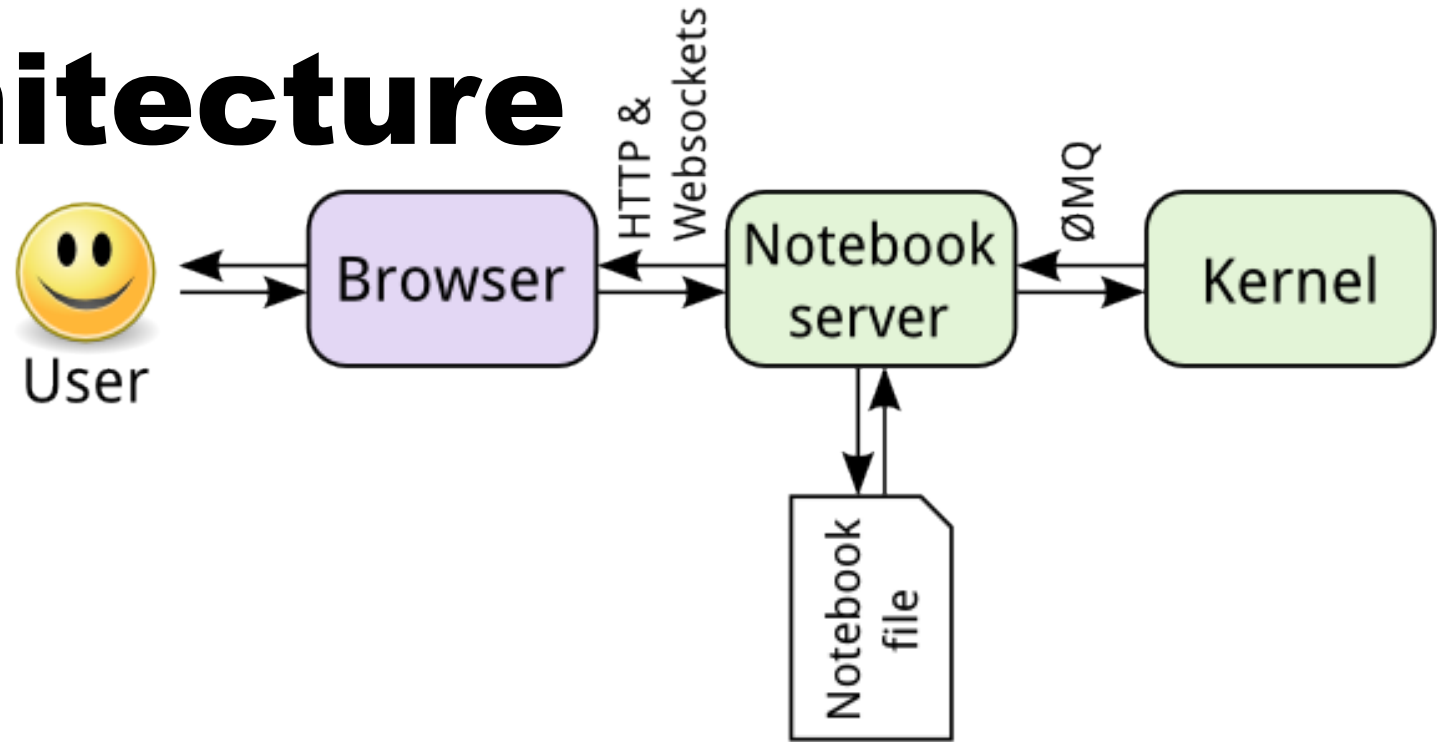
- Jupyter / Docker Architecture
- Using Dockerized Jupyter System
- Building and distributing Docker images
- Virtual Machines & Containers
- Existing Work on Building Docker Images for Jupyter
- Building on Existing Docker Images for Jupyter

# IPython Architecture

- Multiple processes, maybe multi-machine
- Engines evaluate code
  - ipyparallel uses multiple engines
- Controller coordinates communication
- Client provides User Interface
  - Command Line (REPL)
  - GUI (Qt)
  - Web server proxying for browser



# Jupyter Architecture



- Generalization of IPython
  - Language agnostic
  - Kernel evaluates code
  - Web server acts as proxy client
    - Same OS instance as Kernel
  - Browser provides rich multi-media user interface
    - Lots of JavaScript code running on browser


# Detail in Multi-Machine System

If server and kernel on a different machine than the browser:

- Server can up/download notebook file to/from browser
- Persistent data for kernel non-local to browser's host
- Ideally, browser's host and server's host share a filesystem
  - “Big” data should be “near” kernel: high bandwidth & low latency
  - For cloud-hosted server, cloud-storage handy: e.g., Google Drive

# Docker-based Jupyter system

- Docker Container will not have a persistent filesystem
  - Allows Docker image to be updated / upgraded
- Web server started upon container initialization.
- Client must authenticate
  - by server-supplied token
  - server has no way to provide & save user-specific credentials



# Multi-User System

- More advanced, not being covered
- Run JupyterHub on server-side, and it will
  - Manage authentication
  - Start a Docker container for each user
  - Act as proxy between browser and web server in container
  - May start containers on multiple machines to spread out the load

# **Using Dockerized Jupyter**

**What we can do  
once we have  
all the pieces put together**



# Example System

- Windows 10 (Education, Pro, or Enterprise)
  - Hyper-V role installed as optional feature
  - Docker Desktop CE starts Linux VM on system boot
  - Docker containers run in Linux VM
  - Use PowerShell to interact with Docker
- Upcoming WSL2 will be used with future version of Docker
  - WSL2 provides tighter integration of Windows and Linux
  - Linux kernel “inside” Windows 10, still using virtualization HW

# Start-up Procedure (1)

Command Line (Windows PowerShell):

```
docker run --rm -p 8888:8888 -v $HOME/work:/home/joyvan img
```

<b>img</b>	name of Docker filesystem image to start in new container
<b>--rm</b>	remove changes to filesystem image on exit
<b>-p 8888:8888</b>	port map between localhost and container
<b>-v abc:xyz</b>	volume map between user's home and container

Will fail if container already running using port 8888

# Start-up Procedure (2)

- Server starts up in container and prints out URL w/ token
  - URL ends with :8888/
- Start browser
- Copy URL starting with :8888/ from CLI where container started
- In browser's URL entry field:
  - Type `http://localhost`
  - Paste text copied from CLI
  - Should end up with something like

`http://localhost:8888/?token=123abc789def`

# **Start-up Procedure (3)**

- **Browser displays Jupyter start page**
  - **JupyterLab**
    - **Command palette and list of available kernels**
  - **Jupyter Notebook**
    - **Listing of directories & files in mapped directory**

# Misc. Docker Maintenance (1)

To list available images:

```
docker image ls
```

To remove an image:

```
docker rmi foouser/barimage:revision
```

- Can't remove an image that is used by another image
- Can't remove an image used by an active or inactive container

# Misc. Docker Maintenance (2)

To remove all untagged, unused images:

```
docker image prune
```

To remove all inactive containers:

```
docker container prune
```

To list active containers:

```
docker container ls
```

To remove an active container:

```
docker kill foo_bar
```

# Building a Docker file system image (1)

Running Docker-Build is easy:

```
docker build -t paulbuis/cool-jupyter:2019-08-20
```

<b>paulbuis</b>	account name on Docker Hub
<b>cool-jupyter</b>	tag name of image
<b>2019-10-11</b>	revision of tag ( <b>latest</b> if omitted)
	I use YYYY-MM-DD tags

# Building a Docker file system image (2)

Understanding what happens when `docker build` runs is not easy:

- Delayed to later part of talk



# Distributing Image via DockerHub

Place copy of image on `hub.docker.com`:

```
docker push paulbuis/cool-jupyter:2019-08-20
```

will be prompted for credentials

Fetch copy of image from `hub.docker.com`:

```
docker pull paulbuis/cool-jupyter:2019-08-20
```

# **Virtual Machines & Containers**

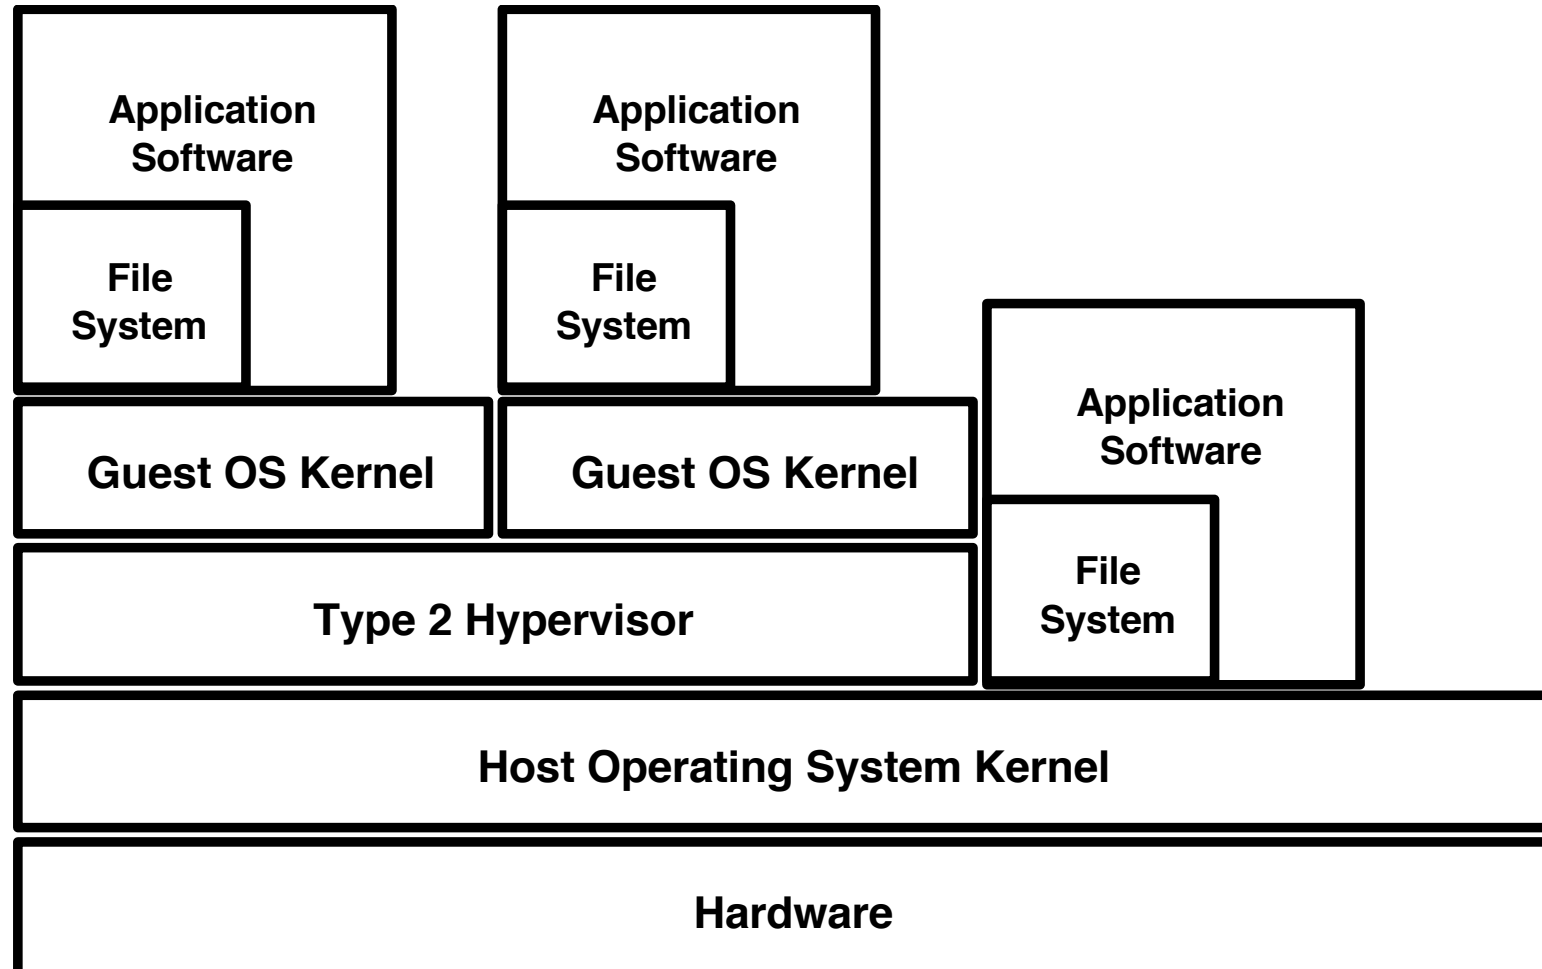
*Heavy & Light  
Process Isolation Techniques*

# Virtual Machines – Type 2

Type 2 Hypervisor  
runs on top of  
Host Operating System


Guest Operating Systems  
runs on top of  
Hypervisor

3 layers of software  
between applications and  
hardware



# Example Type 2 Hypervisors

- Oracle VirtualBox
- Parallels
- VMware
  - Workstation
  - Fusion



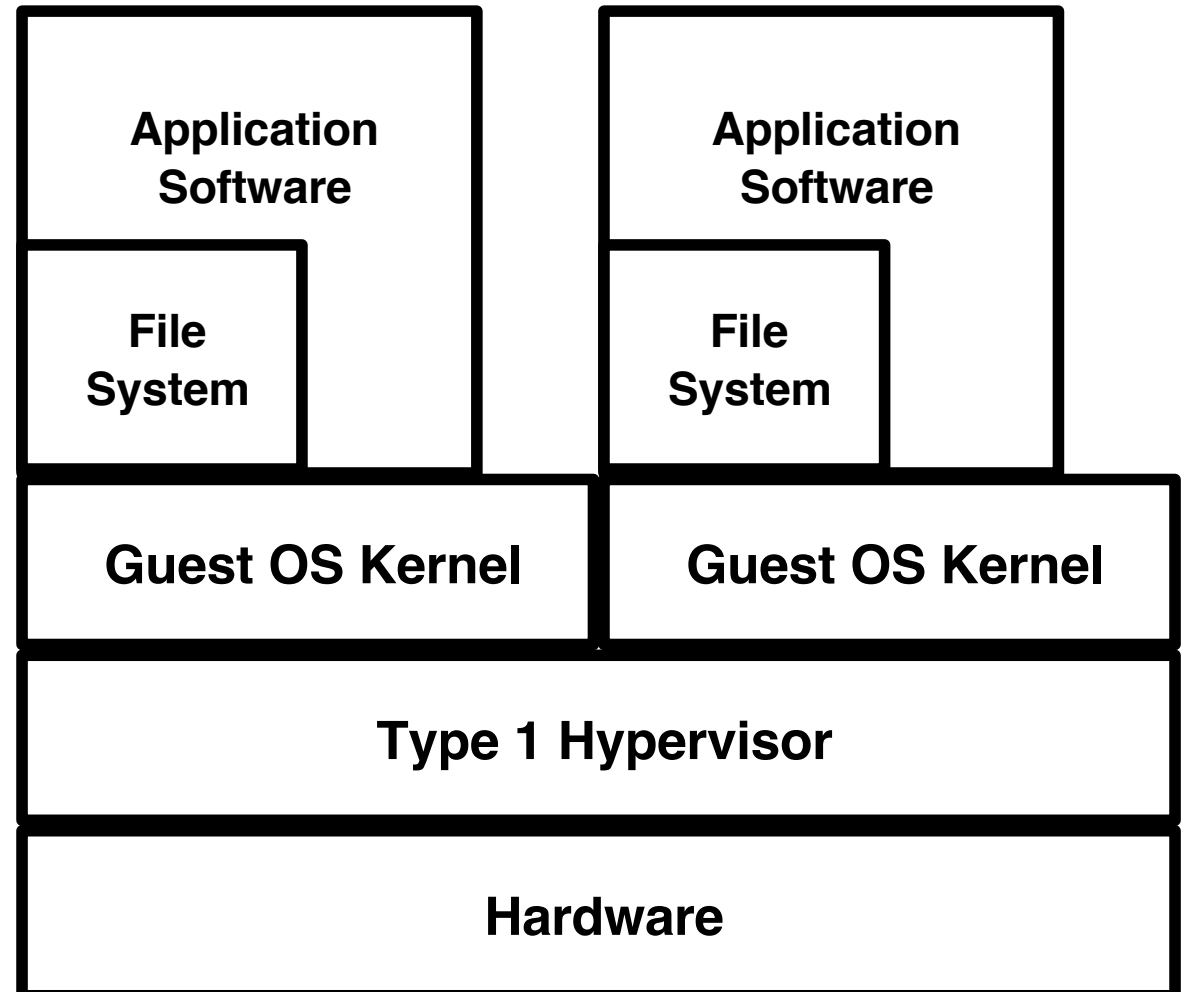
vmware®

# Virtual Machines – Type 1

**Type 1 Hypervisor  
runs on directly on top of  
Hardware**

**Guest Operating Systems  
runs on top of  
Hypervisor**

**2 layers of software  
between applications and  
hardware**



# Example Type 1 Hypervisors

- Microsoft Hyper-V
- Linux KVM
- VMware ESXi




# Containerization

OS provides resource namespaces

Processes limited to resources in a namespace

Container started with limited initd process

1 layer of software between applications and hardware

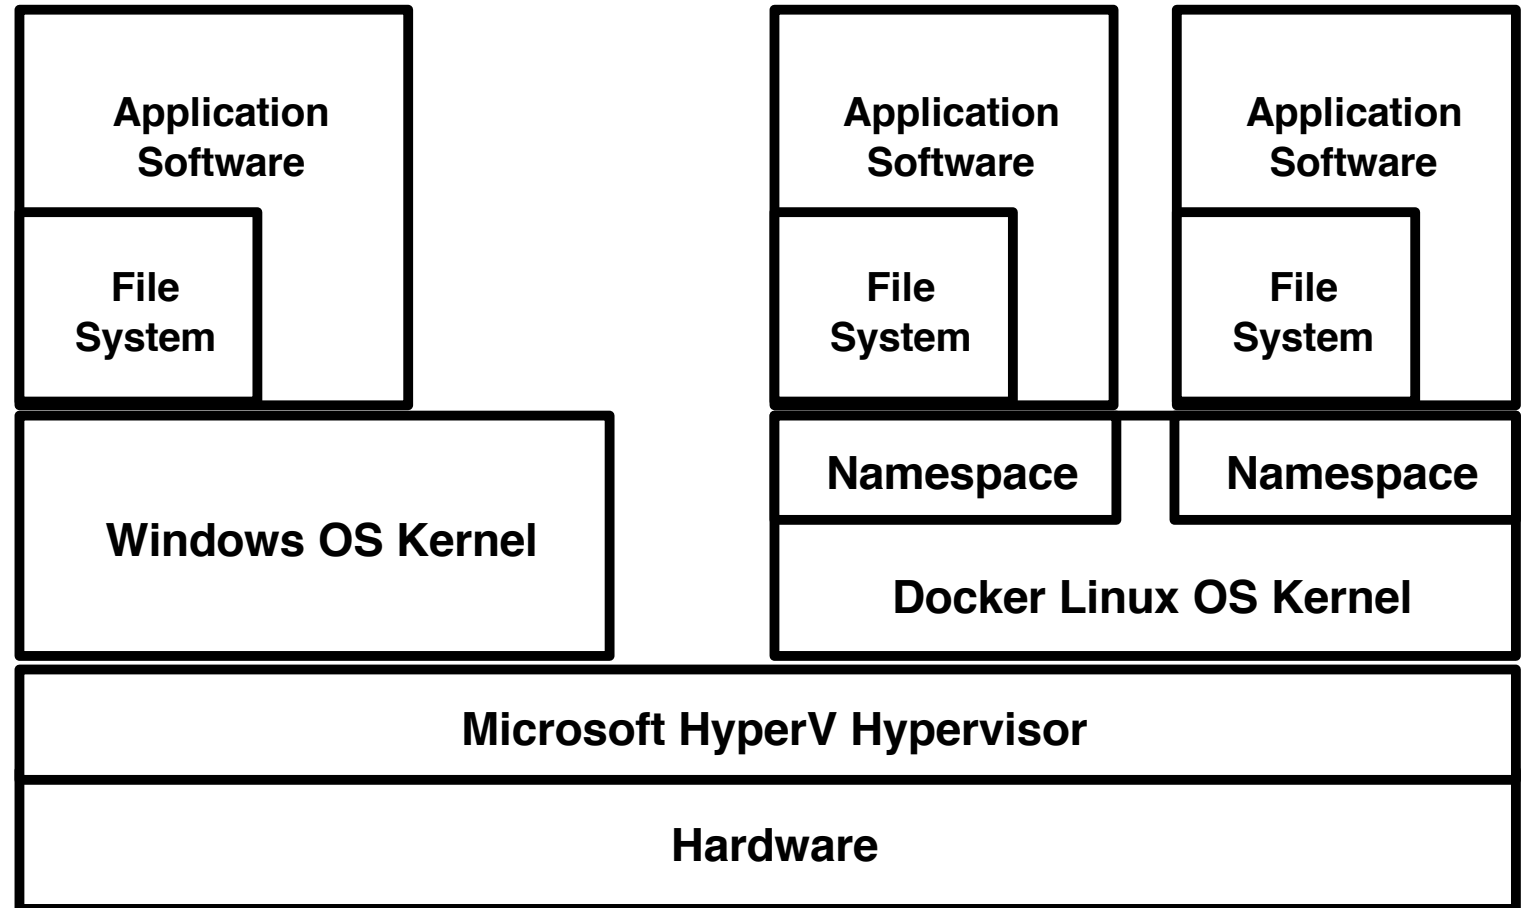


# Docker on Windows (or MacOS) with Linux Containers?

Docker supplied Linux  
OS Kernel runs Docker  
containers

Windows & Linux at  
same layer of stack

2 layers of software  
between applications  
and hardware





# **Another Linux / Hyper-V Alternative**

- **Microsoft provides Linux Distro for Hyper-V, including Ubuntu LTS**
- **Configured to use XRDP with Hyper-V specific client**
- **Feels similar to using Linux on Virtualbox**
- **My desktops & laptops have no problem running Hyper-V with**
  - **Windows 10 Guest**
  - **Docker's Linux Guest running an HTTP server (no UI)**
  - **Microsoft's supplied Ubuntu LTS distro**
  - **All 3 kernels running at the same time: 16 GB RAM, SSD, 2 i7 cores**

# **Jupyter: Architectural Overview**

A Distributed System  
with Standards-based Communication Protocols


# Building Docker Images for Jupyter

Existing GitHub Repositories •



# Building Docker Images: DevOps

- DevOps strives to build computing environments in a reliably reproducible manner using scripting and source code control.
- Used to build platforms on which to deliver commercial software.
- Prevents “It worked on my/our machine, OS & installed software”
- Docker VFS images built from Git repositories with a Dockerfile



Twitter: @paulebuis



git

#PyDataIndy



# The Jupyter Project Provides Base Images

Source:

<https://github.com/jupyter/docker-stacks>

Binary Releases:

<https://hub.docker.com/u/jupyter>

A Docker image is a set of layers forming a VFS

Jupyter Project base image uses:

Official Docker Ubuntu LTS Image

MiniConda Python Distro from Anaconda.org



# Sample Dockerfile Commands

```
FROM ubuntu:bionic-20190612
```

```
USER root
```

```
RUN apt-get update && apt-get -yq dist-upgrade && \  
    apt-get install -yq --no-install-recommends \  
    wget bzip2 ca-certificates sudo locales fonts-liberation
```

```
ENV MINICONDA_VERSION=4.6.14 CONDA_VERSION=4.7.10
```

```
RUN cd /tmp && \  
    wget --quiet https://repo.continuum.io/miniconda/Miniconda3 ...  
    /bin/bash Miniconda3-${MINICONDA_VERSION}-Linux-x86_64.sh -f ...
```

```
...
```

# Sample Dockerfile Commands

...

```
EXPOSE 8888
```

```
ENTRYPOINT ["tini", "-g", "--"]
```

```
CMD ["start-notebook.sh"]
```

```
COPY start-notebook.sh /usr/local/bin/
```

```
USER 1000
```

# Overall Pattern of Jupyter Dockerfiles

- Start with a base image
- Use OS package manager to install packages
- Use `wget` to download other software, then install it
- Use `conda` & `pip` to download and install Python packages
- Use `conda` downloaded tools



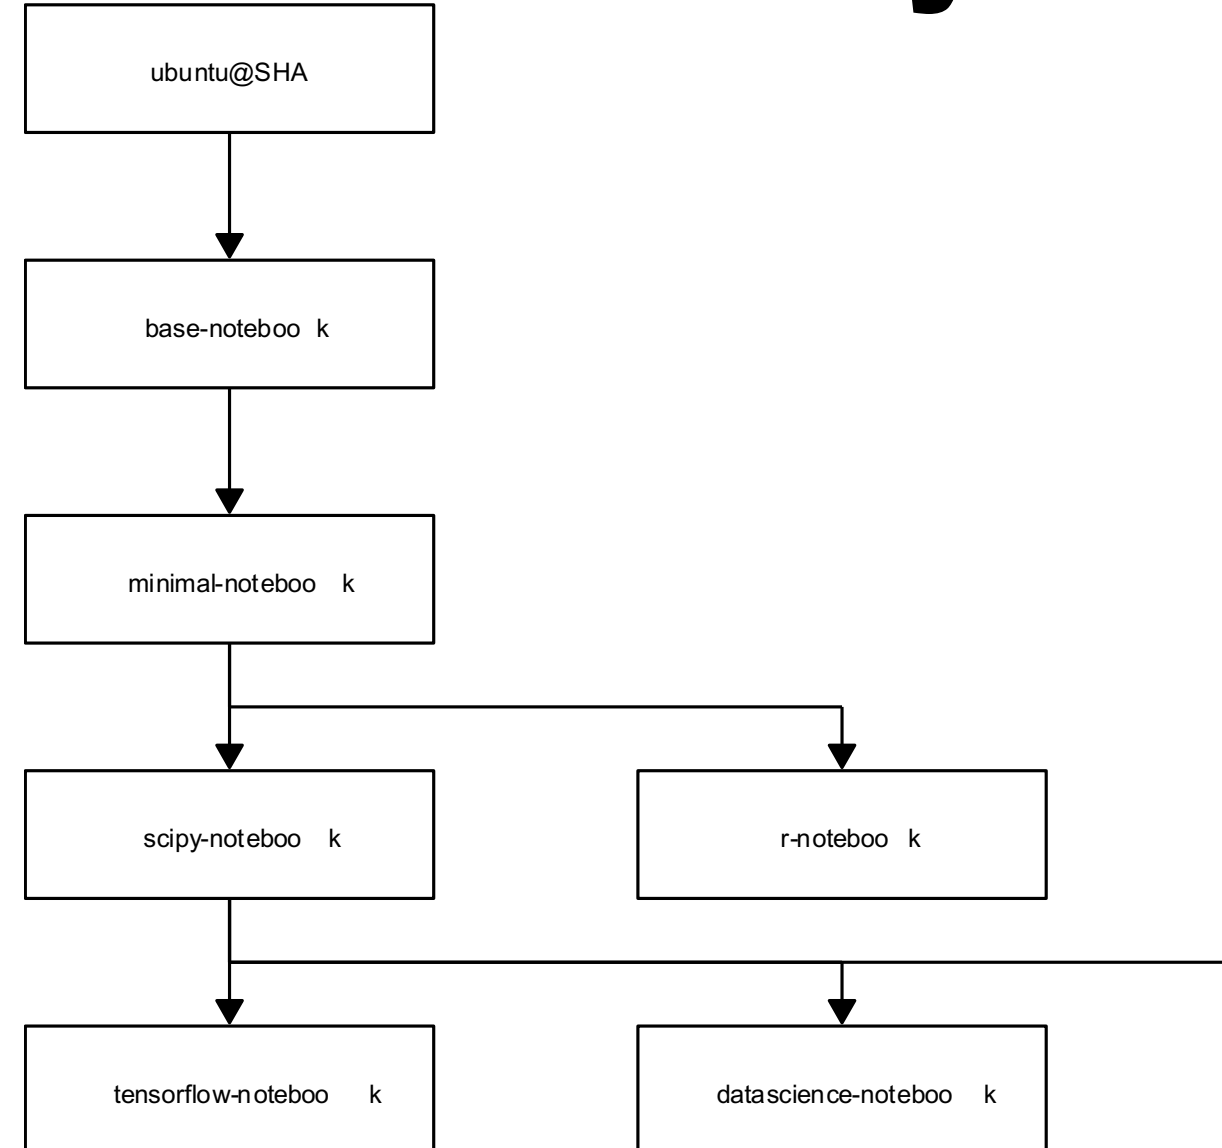
# Jupyter/Docker-Stacks Hierarchy

My GitHub Repository

paulbuis/jupyter-many

Builds on

minimal-notebook



Twitter: @paulebuis

#PyDataIndy